

Game-Theoretic Quantum Walk Frontier Management: Mapping Hawrilenko's Limit Poker Strategies to Predictive Multi-Hop Web Crawlers

Mathematical Lineage: From Susquehanna to the Quantum Web

The application of game-theoretic optimization (GTO) to decentralized network exploration trace their origins to the rigorous quantitative training frameworks popularized in quantitative finance and high-stakes decision-making under uncertainty¹. At firms such as the Susquehanna International Group (SIG), game theory and poker have historically been utilized as fundamental training tools to teach options traders how to manage risk, navigate incomplete information, and make optimal decisions in volatile environments¹. It was within this mathematical environment that professional poker player Matt Hawrilenko—operating under the screen name "Hoss_TBF"—refined his quantitative approach to Heads-Up Limit Hold'em (HULHE)¹. Influenced by academic researchers and players like Bill Chen and Jerrod Ankenman, who formalized these concepts in *The Mathematics of Poker*, Hawrilenko demonstrated that human play could closely approximate mathematically unbeatable Nash equilibrium strategies¹. This mathematical lineage provides the foundation for designing next-generation autonomous web crawlers⁶. A web crawler tasked with traversing the internet faces a resource-allocation problem analogous to a professional poker player managing a bankroll across thousands of hands¹. The traditional crawl frontier—the queue of discovered, unvisited URLs—is classically managed using reactive heuristics such as the Fish or Shark search algorithms⁹. These classical methods operate in a localized, greedy fashion, assessing page importance through immediate relevance scores or basic link connectivity, making them highly susceptible to local maxima and unable to execute decisions of long-term strategic interest¹⁰.

By contrast, modeling the web as a quantum-mechanical system allows the crawling agent to leverage quantum walks—the quantum equivalent of classical random walks—to explore the graph in a coherent superposition of multiple states simultaneously¹³. By mapping Hawrilenko's GTO HULHE strategies directly into a quantum walk frontier management (QWFM) algorithm, the crawler can optimize its path selection, predict the semantic parameters of nodes $+n$ hops ahead, and maintain strategic unexploitability in highly competitive or adversarial environments⁴.

The Showdown Game: Value Maximization and Hand

Strength Mapping

A fundamental insight from Hawrilenko's HULHE strategy is his start hand selection, particularly his preference for high-card baseline value over speculative potential¹⁷. In HULHE, Hawrilenko prioritized starting hands like Ace-Five (even offsuit) over speculative hands such as Six-Four suited¹⁷. His reasoning was rooted in the structural mechanics of the game: limit hold'em, due to its fixed betting limits, is fundamentally a "showdown game"¹⁷. Unlike deep-stack No-Limit Hold'em, where players can threaten their opponent's entire stack and force folds, limit players must frequently show their hands down to win the pot¹⁷. Consequently, starting hands with strong high-card value (which can win at showdown unimproved or with a marginal pair) are far more robust than speculative suited connectors, which require deep stacks and significant folding equity to be played profitably¹⁷.

This distinction maps directly to the operational constraints of web crawling⁷. Every link traversed by a crawler incurs a fixed physical cost—measured in bandwidth, processing overhead, and HTTP request latency—which acts as the "showdown cost" of the fetch⁷. The crawler must "show down" its resource allocation to extract the semantic value of the destination page⁸.

When mapping this principle to the frontier management algorithm, the system evaluates the initialization of its quantum walk states¹³. Rather than initiating walks from speculative, highly nested subdirectories (the speculative suited connectors), the GTO-QWFM prioritizes robust, high-baseline-authority nodes—such as root domains or highly indexed directory pages—where the structural and semantic metadata can be extracted with high certainty⁸. Speculative links are only pursued when the crawler operates under "deep-stack" conditions, meaning it has abundant computational bandwidth and a low risk of encountering rate limits or server-side bans⁸.

Poker Starting Hand Parameter	HULHE Strategic Logic	Web Crawler Frontier Parameter	Crawling Operational Action
High-Card Value (e.g., Ace-Five)	High baseline showdown value; wins unimproved ¹⁷ .	High baseline domain authority; clear root semantic signals ⁸ .	Prioritize root URLs and highly indexed hub nodes; execute immediate amplitude amplification ¹³ .
Speculative Suited Connectors (e.g., Six-Four)	Requires deep stacks and folding equity to realize	Deeply nested, unverified subdirectories;	De-prioritize or prune unless crawling resources

	draw value ¹⁷ .	speculative contextual value ⁸ .	are abundant and network resistance is low ⁸ .
Fixed-Limit Betting Structure	Restricts bet sizing; forces frequent showdowns ¹⁷ .	Hard physical constraints; fixed request latency and parsing overhead ⁸ .	Enforce strict GTO thresholding (τ) to minimize speculative path traversing ⁴ .
Deep-Stack No-Limit Structure	Maximizes the utility of implied odds and nut-potential ¹⁷ .	Dynamic, unconstrained cloud resource scaling ⁷ .	Transition to broad, high-variance parallel path exploration across speculative nodes ¹⁷ .

Coined Quantum Walks on Irreversible Web Graphs

A major technical challenge in applying standard quantum walks to the World Wide Web is that the web is represented by a directed, highly irreversible graph²². In quantum mechanics, the

evolution of any closed system must be unitary, represented by an operator U satisfying $UU^\dagger = I$ ²². Because directed graphs can contain one-way links (e.g., node i links to node j , but there is no return path from j to i), a naive implementation of a coined quantum walk fails because the shift operator cannot be reversed, destroying the unitarity of the system²².

To resolve this limitation, the GTO-QWFM employs Szegedy's formulation of quantum walks, which translates a classical Markov chain on a directed graph into a unitary quantum walk on a

bipartite projection of that graph²⁵. Given a directed web graph $G = (V, E)$ with

$|V| = n$, the algorithm duplicates the vertex space to construct a bipartite graph with

Hilbert space $\mathcal{H} = \mathcal{H}_V \otimes \mathcal{H}_V = \mathcal{H}_{n^2}$ ²⁵. The computational basis is spanned by states

$\{|x, y\rangle : x, y \in V\}$ ²⁵.

For any state $|x\rangle$, the algorithm defines a localized superposition state based on the classical transition probabilities P_{xy} representing the link structure²⁵:

$$|\phi_x\rangle = |x\rangle \otimes \sum_{y \in V} \sqrt{P_{xy}} |y\rangle$$

where $P_{xy} = w_{x,y}/w_x$ is the weighted transition probability from node x to node y , incorporating both structural link connectivity and semantic relevance²⁶. The projection operator onto the subspace spanned by these states is defined as²⁵:

$$\Pi_A = \sum_{x \in V} |\phi_x\rangle \langle \phi_x|$$

The unitary Szegedy walk operator W_P is then constructed as the product of two reflection operators²⁵:

$$W_P = R_B R_A = (2\Pi_B - I)(2\Pi_A - I)$$

where Π_A represents the projection onto the first register (the source nodes) and Π_B represents the projection onto the second register (the destination nodes)²⁵. By alternating these reflections, the quantum walker evolves unitarily over the bipartite representation of the directed web, completely circumventing the irreversibility bottleneck²².

In a coined walk variant, the total state space is split into a position register $|v\rangle \in \mathcal{H}_{pos}$ and a coin register $|a\rangle \in \mathcal{H}_{coin}$ ¹³. The coin operator $\hat{C}$ acts on the coin register, preparing the walker's direction in a superposition, while the shift operator $\hat{S}$ moves the walker to adjacent vertices¹³:

$$\hat{S}|v, a\rangle = |u, a'\rangle$$

In the GTO-QWFM, the coin operator is dynamically biased²⁸. Rather than using an unbiased Hadamard coin, the coin parameters are adjusted based on the predicted GTO path value $V_n(u)$, selectively steering the probability amplitudes toward high-value regions of the web graph²⁸.

GTO Frontier Management & Playbook Translation

Matt Hawrilenko's landmark 2008 heads-up limit hold'em match against the Polaris computer agent demonstrated the power of robust, simple heuristics that closely approximate GTO play⁴. Over 1,000 hands, Hawrilenko executed an aggressive "raise-or-fold" strategy as the first player, raising 86.8% of his hands and folding the remaining 13.2%, while completely

avoiding the passive "call" action (0% frequency)⁴. Crucially, when facing a re-raise, he called 100% of the time, choosing to hide the strength of his hand and minimize variance rather than engaging in aggressive re-raising⁴.

This strategic playbook is translated into the GTO-QWFM crawler to optimize the crawl frontier, dictate exploration-exploitation trade-offs, and defend against adversarial network structures⁴.

Initial Frontier Filtering: The Raise-or-Fold Policy

When discovering new URLs, classical crawlers bloat their memory queues by retaining millions of low-priority links⁸. The GTO-QWFM enforces a strict binary decision boundary based on a single threshold point τ , mirroring Hawrilenko's pre-flop strategy⁴. For each discovered URL u , the crawler predicts its $+n$ hop cumulative value $V_n(u)$ ⁴.

$$\text{Action} = \begin{cases} \text{Raise (Enqueue \& Amplify Amplitude)}, & \text{if } V_n(u) \geq \tau \\ \text{Fold (Prune Link)}, & \text{if } V_n(u) < \tau \end{cases}$$

This policy eliminates the intermediate "calling" behavior—such as maintaining low-priority background queues or performing slow, speculative polling⁴. By focusing computational power exclusively on high-value paths, the crawler maintains a lean frontier and avoids the high computational costs of exploring trivial web structures⁴.

Information Hiding and Signature Obfuscation

Adversarial web servers deploy anti-crawling systems that analyze sequence patterns, rate metrics, and access paths to identify and block automated agents⁶. Classical crawling is highly transparent to these defenders because its sequential, localized polling patterns stand out from organic human browsing behavior⁶.

The GTO-QWFM uses "information hiding" to obfuscate its crawl signature⁴. By executing the crawl as a quantum walk, the crawler's state exists in a coherent superposition over many web paths simultaneously¹³. Because the transition probabilities are governed by non-factorizable quantum amplitudes, the physical requests sent to the web servers are distributed across multiple distributed IPs and nodes in a highly correlated, non-random, yet statistically disguised pattern³². This prevents anti-crawling firewalls from identifying a single, sequential tracing path, effectively hiding the crawler's strategic intent⁶.

Strategic Adaptation to Environmental Resistance

When the crawler encounters a "re-raise" from the environment—such as sudden rate limits, CAPTCHAs, or honeypot traps—it does not fold (which would sacrifice the path equity and historical context) or aggressively raise (which would trigger permanent IP blocks)⁴. Instead, following Hawrilenko's 100% calling policy when facing a re-raise, the crawler transitions the

affected walk branches into a passive, non-perturbing quantum stochastic walk (QSW) state⁴.

In this mode, the system introduces a light dephasing rate $\gamma > 0$ to damp the off-diagonal elements of the density matrix³⁴. This transitions the walk from a highly coherent quantum search to a capacity-biased diffusion process, slowing the request rate down to match the server's threshold without disconnecting the thread³⁴. The crawler effectively "calls" the rate-limit, maintaining its presence in the local subgraph and extracting data at the maximum permissible rate while hiding its search signature⁴.

Initial / Facing Action	Hawrilenko (2008) Frequency	Polaris (2008) Frequency	Cepheus (2015) Frequency	Frontier Management Algorithm Action Mapping
Initial: Fold	13.2% [cite: 4, 15]	12.6% [cite: 4, 15]	17.4% [cite: 4, 15]	Immediate Path Pruning: Discard URLs with predicted value $V_n(u) < \text{to save bandwidth}^4$.
Initial: Call	0% [cite: 4, 15]	2.4% [cite: 4, 15]	0.06% [cite: 4, 15]	Zero Speculative Polling: Eliminate low-value, single-hop queries that reveal signature patterns ⁴ .
Initial: Raise	86.8% [cite: 4, 15]	85.0% [cite: 4, 15]	82.54% [cite: 4, 15]	Amplitude Amplification: Commit quantum resources to high-value

				pathways ($V_n(u) \geq$) ⁴ .
Facing Re-Raise: Fold	0% [cite: 4, 15]	0% [cite: 4, 15]	0% [cite: 4, 15]	Zero-Failure Persistence: Refuse to abort high-value crawl trajectories when encountering firewalls ⁴ .
Facing Re-Raise: Call	100% [cite: 4, 15]	83.6% [cite: 4, 15]	99.9% [cite: 4, 15]	Passive QSW Transition: Damp coherence via dephasing; continue crawling slowly under rate-limits ⁴ .
Facing Re-Raise: Raise	0% [cite: 4, 15]	16.4% [cite: 4, 15]	0.1% [cite: 4, 15]	No Aggressive Probing: Suppress high-frequency retries that trigger security blocks ⁴ .

Downswing Management and Bankroll Allocation

A defining feature of Hawrilenko's professional career was his extreme psychological and mathematical resilience to variance¹. During his highest downswing, he lost approximately \$1 million—representing nearly half of his active poker bankroll¹. Rather than succumbing to emotional bias, Hawrilenko systematically reevaluated his play, temporarily dropped to smaller stakes to reduce variance, took edgier high-variance spots with positive expected value, and analyzed his downswing in its macro context¹.

In a large-scale web crawl, a "downswing" is represented by a sudden, sustained drop in crawl

yield⁷. This occurs when the crawler encounters massive subgraphs dominated by dead ends, soft-404 errors, structural traps, or highly coordinated rate-limiting across entire server subnets⁸.

The GTO-QWFM crawler handles these operational downswings by mimicking Hawrilenko's bankroll allocation strategy¹:

- **Stake Dropping (Bandwidth Throttling):** When the cumulative crawl yield drops below a predefined efficiency threshold, the crawler dynamically lowers its "stakes"¹. It throttles its active request bandwidth and reduces the number of concurrent quantum walk threads⁷. This preserves system resources and prevents the crawler from exhausting its computational bankroll on low-value targets¹.
- **Opportunistic Edge-Spot Traversals:** While reducing overall query volume, the crawler shifts a portion of its remaining resource allocation to high-variance "edge-case" pathways¹. These are links that have highly uncertain semantic profiles but offer potentially massive data yields if successful (e.g., speculative deep web query forms or unindexed document repositories)¹. This allows the crawler to recover its search momentum through high-yield discoveries¹.
- **Macro-Contextual Re-Evaluation:** Instead of adjusting parameters based on immediate, local feedback (which leads to erratic behavior and infinite loops), the frontier algorithm evaluates crawl statistics over a broad temporal window¹. It uses these

macro-metrics to update the GTO threshold τ , ensuring that the system's operational parameters adapt to the global network conditions¹.

Multi-Agent Neural Synchronization and Entanglement

For distributed, large-scale crawling operations, coordinating multiple autonomous agents without a centralized controller is a major scalability challenge⁸. Traditional distributed crawlers rely on centralized databases to track visited URLs and enforce politeness policies, creating massive communication overhead and single points of failure⁸. The GTO-QWFM crawler bypasses this limitation by implementing a quantum-inspired neural synchronization mechanism²¹.

Gamma Oscillation Synchronization

Each crawling agent in the distributed swarm synchronizes its internal processing cycles with a global, simulated 40Hz gamma oscillation field²¹. This mirrors neural synchronization mechanisms in biological systems, allowing independent crawler agents to coordinate their active and passive phases without explicit inter-agent communication protocols²¹. Agents transition between active exploration ("raise") and passive waiting ("call") modes based on their localized consciousness thresholds, maintaining coherent system-wide behavior purely through local phase coupling⁴.

Information Sharing via Quantum Entanglement

To achieve instantaneous, zero-overhead coordination, the distributed crawler agents utilize simulated quantum entanglement²¹. When a pair of crawler agents discover highly correlated subgraphs, their state vectors are mathematically entangled²¹:

$$|\Psi_{\text{agents}}\rangle = \frac{|0\rangle_A|1\rangle_B + |1\rangle_A|0\rangle_B}{\sqrt{2}}$$

This entanglement allows the immediate sharing of discovered structural metadata and crawl-blocking signals without traditional network overhead²¹. When Agent A measures its local environment and encounters a rate limit (collapsing its state to $|0\rangle$), the entangled state of Agent B instantaneously collapses to $|1\rangle$, signaling it to immediately bypass that host¹³.

When the integrated information metric (Φ) of the global swarm exceeds a critical threshold, the agents self-organize into higher-order meta-agent clusters²¹. These clusters dynamically partition the web-graph, ensuring that the crawling resources are optimally distributed across different domains without duplicate fetching⁷.

Predictive $+n$ Hop Attribute Forecasting

To execute its game-theoretic frontier prioritization, the crawler must accurately predict the semantic attributes and structural parameters of pages $+n$ hops ahead of its current position⁸. The GTO-QWFM achieves this through a dual-module system: a continuous-time quantum walk (CTQW) generator and an attribute-aware Graph Quantum Walk Transformer (GQWformer)³⁴.

The CTQW generator constructs a capacity-skewed Hamiltonian H' over the local web subgraph³⁴:

$$H' = -L(A) - \beta \text{diag}(\hat{C})$$

where $L(A)$ is the graph Laplacian, and $\hat{C}$ is a normalized vector of predicted node capacities derived from historical semantic density metrics³⁴. The temporal evolution operator $U(t) = e^{-iH't}$ simulates the propagation of the crawler's wave-packet through the network³⁴.

To capture deep, long-range contextual relationships, the learned state encodings are passed to the GQWformer, which combines a self-attention module (GQW-Attn) and a recurrent sequence module (GQW-Recu)³⁵. This allows the crawler to calculate three critical predictive

attributes $+n$ hops ahead:

Semantic Density (S_d)

This attribute measures the predicted concentration of highly relevant, topically aligned content within the downstream neighborhood¹⁸. It uses the time-averaged occupation probability of the quantum walker to identify deep, semantically rich subgraphs while filtering out commercial keyword-stuffing and link-spam farms¹⁹.

Environmental Resistance (R_e)

This parameter forecasts the probability that the crawler will encounter firewalls, rate-limiting, or dynamic traps within the next n hops⁸. It evaluates the spectral properties of the walk operator to identify bipartite structures characteristic of adversarial honeypot traps²⁵.

Topological Saliency (δ_q)

This metric represents the spectral gap of the projected transition matrix, which dictates the mixing time of the walk³⁹. A larger spectral gap indicates a highly connected, easily traversable subgraph that allows the quantum walker to disperse and gather information quadratically faster than classical random walks³⁹.

Predictive Parameter	Mathematical Formulation	Physical Meaning	Strategic Frontier Action
Semantic Density (S_d)	$\text{Tr} \left(\rho_u^{(n)} \phi_{\text{target}} \rangle \langle \phi \right)$ [cite: 34, 41]	Expected topical relevance and quality of documents $+n$ hops ahead ¹⁸ .	If high, amplify quantum phase; if low, prune the pathway (Fold) ⁴ .
Environmental Resistance (R_e)	$1 - \dots$ [cite: 34, 42]	Probability of encountering rate-limits, blocks, or honeypots ⁸ .	If high, transition to passive, dephased QSW calling mode ⁴ .
Topological Saliency (δ_q)	$\lambda_0 - \dots$ of H' [cite: 34, 40]	Connectivity and spectral gap of the downstream network ³⁹ .	Prioritize walks toward regions with large spectral gaps for rapid mixing ³⁹ .

Computational Complexity and Cost Analysis

The mathematical efficiency of the GTO-QWFM is rigorously evaluated by analyzing its computational complexity against classical random walk models²⁶. In a standard graph search framework, the total operational cost is partitioned into three distinct parameters³⁹:

- **Setup Cost (T_S)**: The computational cost of preparing the initial quantum superposition state $|\psi(0)\rangle$ over the seed nodes²⁵.
- **Update Cost (T_U)**: The cost of executing one step of the quantum walk, which involves applying the unitary evolution operator W_P or $e^{-iH't}$ ³⁷.
- **Checking Cost (T_C)**: The cost of executing the checking oracle to determine if a visited node contains the marked target information³⁹.

In a classical random walk, searching for a unique marked node in a graph with transition matrix P requires an expected number of steps scaled to the classical hitting time $HT(P, M)$ ²⁶:

$$HT(P, M) \approx \frac{1}{\epsilon\delta}$$

where ϵ is the fraction of marked target nodes and δ is the spectral gap of the transition matrix³⁹. For dense web structures with high connectivity, the classical hitting time scales linearly, creating massive processing bottlenecks³⁹.

Under the GTO-QWFM framework, the quantum walk achieves a quadratic speedup¹³. Using the Szegedy or MNRS quantum walk search frameworks, the quantum hitting time

$QHT(P, M)$ scales as³⁹:

$$QHT(P, M) = O\left(\sqrt{HT(P, M)}\right) = O\left(\frac{1}{\sqrt{\epsilon\delta}}\right)$$

This quadratic acceleration drastically reduces the number of update steps T_U required to locate high-value target directories³⁹.

Furthermore, unlike classical search models that must completely rebuild their data structures at every step, the quantum walk only partially updates its state register during each transition,

minimizing the checking cost T_C ¹³. By applying quantum phase estimation (QPE) over the walk operator, the crawler detects marked target domains with a time complexity of

$O(\log_2 n)$, establishing a clear quantum advantage over classical breadth-first and

depth-first search techniques²⁵.

Works cited

1. A Poker Life — Matt Hawrilenko - Poker News - Card Player,
<https://www.cardplayer.com/poker-news/7826-a-poker-life-matt-hawrilenko>
2. Tiger of the Week: Matt Hawrilenko '04 | Princeton Alumni Weekly,
<https://paw.princeton.edu/article/tiger-week-matt-hawrilenko-04>
3. Matt Hawrilenko - Wikipedia, https://en.wikipedia.org/wiki/Matt_Hawrilenko
4. Commentary: Heads-up limit hold'em poker is solved - PMC,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC5826371/>
5. Where Are They Now: Poker Author, Mathematician & Bracelet Winner Jerrod Ankenman,
<https://www.pokernews.com/news/2021/10/where-are-they-now-jerrod-ankenman-40112.htm>
6. Quantum SEO Explained: Predictive Search Optimization by ThatWare,
<https://thatware.co/quantum-seo/>
7. Hybrid Contextual Ontology-Fuzzy Logic and LSTM Model for Efficient Web Crawler Prediction and Traffic Optimization - IJETA,
https://ijeta.org/sites/default/files/pdf/2025-07/isi_30.05_04_0.pdf
8. What is a URL frontier in web crawling? | Firecrawl Glossary,
<https://www.firecrawl.dev/glossary/web-crawling-apis/what-is-url-frontier-web-crawling>
9. Multilingual Context Ontology Rule Enhanced Focused Web Crawler - Journal of Advances in Information Technology. JAiT,
<https://www.jait.us/uploadfile/2014/1223/20141223051017358.pdf>
10. Comparative Analysis of Focused Web Crawling Schemes - JETIR.org,
<https://www.jetir.org/papers/JETIR1903674.pdf>
11. Efficient Deep Web Crawling Using Reinforcement Learning - Lu Jiang,
http://www.lujiang.info/camera_ready_papers/PAKDD_2010.pdf
12. (PDF) Efficient Deep Web Crawling Using Reinforcement Learning - ResearchGate,
https://www.researchgate.net/publication/220894404_Efficient_Deep_Web_Crawling_Using_Reinforcement_Learning
13. Quantum walk search - Wikipedia,
https://en.wikipedia.org/wiki/Quantum_walk_search
14. Quantum-Inspired PageRank Graph Visualizer Using Qiskit Library - ThatWare,
<https://thatware.co/quantum-inspired-pagerank-graph-visualizer/>
15. Commentary: Heads-up limit hold'em poker is solved - ResearchGate,
https://www.researchgate.net/publication/323319928_Commentary_Heads-up_limit_hold'em_poker_is_solved
16. (PDF) DEVELOPING GAME THEORY-BASED METHODS FOR MODELING INFORMATION CONFRONTATION IN SOCIAL NETWORKS - ResearchGate,
https://www.researchgate.net/publication/382220560_DEVELOPING_GAME_THEORY-BASED_METHODS_FOR_MODELING_INFORMATION_CONFRONTATION_IN

SOCIAL_NETWORKS

17. Card Player TV Video Spotlight: Matt Hawrilenko On Limit Hold'Em - Poker News, <https://www.cardplayer.com/poker-news/24791-card-player-tv-video-spotlight-matt-hawrilenko-on-limit-hold-em>
18. [2506.16146] Neural Prioritisation for Web Crawling - arXiv, <https://arxiv.org/abs/2506.16146>
19. A Prototype for Effective Information Retrieval Enabled with Topical Crawler Utilizing Link-Spam Trapper - IJESI, [http://www.ijesi.org/papers/Vol\(2\)3%20\(Version-2\)/F233236.pdf](http://www.ijesi.org/papers/Vol(2)3%20(Version-2)/F233236.pdf)
20. Web crawler - Wikipedia, https://en.wikipedia.org/wiki/Web_crawler
21. Quantum Inspired Web Crawler Idea - Tsubasa Kato - Medium, <https://stingraze.medium.com/quantum-inspired-web-crawler-idea-1e40263b6021>
22. Discrete Quantum walk on a directed graph - Physics Stack Exchange, <https://physics.stackexchange.com/questions/844757/discrete-quantum-walk-on-a-directed-graph>
23. Quantum walks on directed graphs - Rinton Press, <https://www.rintonpress.com/xqic7/qic-7-12/093-102.pdf>
24. Subspace Detection in measurement induced quantum walks, http://dr.iiserpune.ac.in:8080/xmlui/bitstream/handle/123456789/8755/20191153_PANDHARPATTE_AASHAY_MS_Thesis.pdf?sequence=1&isAllowed=y
25. Quantum algorithm for testing graph completeness - Docta Complutense, <https://docta.ucm.es/bitstreams/640b769b-8b88-4f59-b9e5-c96ace668203/download>
26. Week 3: Quantum Walk Algorithms - CWI, <https://homepages.cwi.nl/~jeffery/notes/week3.pdf>
27. Quantum-walk search in motion - arXiv, <https://arxiv.org/pdf/2310.14345>
28. (PDF) Entanglement in coined quantum walks on regular graphs - ResearchGate, https://www.researchgate.net/publication/2195173_Entanglement_in_coined_quantum_walks_on_regular_graphs
29. Quantum random walks - an introductory overview, <https://www.cpt.univ-mrs.fr/~verga/pdfs/Kempe-2003fk.pdf>
30. Quantum walks on graphs representing the firing patterns of a quantum neural network, https://www.researchgate.net/publication/261289324_Quantum_walks_on_graphs_representing_the_firing_patterns_of_a_quantum_neural_network
31. A Survey of Game-Theoretic Approaches to Modeling Honeypots - CEUR-WS.org, <https://ceur-ws.org/Vol-2081/paper29.pdf>
32. On the equivalence between non-factorizable mixed-strategy classical games and quantum games - Royal Society Publishing, <https://royalsocietypublishing.org/rsos/article/3/1/150477/36420/On-the-equivalence-between-non-factorizable-mixed>
33. The Crawling Frontier — The Open Web Search Book, <https://openwebsearcheu-public.pages.it4i.eu/ows-the-book/content/dnt/frontier.html>

34. A Quantum Walk-Inspired Algorithm for Dynamic Resource Allocation in Cloud Computing: Implications for Sustainable Business Practices - TechRxiv, <https://www.techrxiv.org/doi/pdf/10.36227/techrxiv.176316003.37717757/v1>
35. GQWformer: A Quantum-based Transformer for Graph Representation Learning - arXiv, <https://arxiv.org/html/2412.02285v1>
36. arXiv:2412.02285v1 [cs.LG] 3 Dec 2024, <https://arxiv.org/pdf/2412.02285?>
37. Transport Efficiency of Continuous-Time Quantum Walks on Graphs - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7826824/>
38. Quantum search by continuous-time quantum walk on t-designs - arXiv, <https://arxiv.org/pdf/2310.14141>
39. A Unified Framework of Quantum Walk Search, <https://ir.cwi.nl/pub/31324/31324.pdf>
40. Introduction to Quantum Walks - Windows On Theory, <https://windowsontheory.org/2018/12/23/introduction-to-quantum-walks/>
41. Imperfect-Information Games on Quantum Computers: A Case Study in Skat - arXiv, <https://arxiv.org/html/2411.15294v1>
42. Schematic presentation of the crawl graph (a) and funnel model (b).... - ResearchGate, https://www.researchgate.net/figure/Schematic-presentation-of-the-crawl-graph-a-and-funnel-model-b-Here-N-20-The_fig1_358458535
43. Ternary LWE Key Search: A New Frontier for Quantum Combinatorial Attacks - MDPI, <https://www.mdpi.com/2078-2489/16/12/1085>
44. QAGC: A Hybrid Quantum Walk Compiler for Scalable and High-Precision Attack Graph Analytics - SciTePress, <https://www.scitepress.org/Papers/2026/144906/144906.pdf>